

# Homework 2 - Berkeley STAT 157

Handout 1/29/2019, due 2/5/2019 by 4pm in Git by committing to your repository.

```
In [1]: from mxnet import nd, autograd, gluon
        from matplotlib import pyplot as plt
        import numpy as np
```

## 1. Multinomial Sampling

Implement a sampler from a discrete distribution from scratch, mimicking the function `mxnet.ndarray.random.multinomial`. Its arguments should be a vector of probabilities  $p$ . You can assume that the probabilities are normalized, i.e. that they sum up to 1. Make the call signature as follows:

```
samples = sampler(probs, shape)
```

```
probs    : An ndarray vector of size n of nonnegative numbers summing up to 1
shape    : A list of dimensions for the output
samples  : Samples from probs with shape matching shape
```

Hints:

1. Use `mxnet.ndarray.random.uniform` to get a sample from  $U[0, 1]$ .
2. You can simplify things for probs by computing the cumulative sum over probs.

```
In [2]: def sampler(probs, shape):
        ## Add your codes here
        samples = nd.zeros(shape).reshape(-1,)
        num_items = shape[0] * shape[1]
        for i in range(num_items):
            ran_gen = nd.random.uniform()
            counter = 0
            while ran_gen > 0:
                ran_gen -= probs[counter]
                if (ran_gen) < 0:
                    break
                counter += 1
            samples[i] = probs[counter]
        samples = samples.reshape(shape)
        return samples

# a simple test
sampler(nd.array([0.2, 0.3, 0.5]), (2,3))
```

```
Out[2]: [[0.5 0.5 0.5]
         [0.5 0.5 0.5]]
<NDArray 2x3 @cpu(0)>
```

## 2. Central Limit Theorem

Let's explore the Central Limit Theorem when applied to text processing.

- Download <https://www.gutenberg.org/ebooks/84> (<https://www.gutenberg.org/files/84/84-0.txt>) from Project Gutenberg
- Remove punctuation, uppercase / lowercase, and split the text up into individual tokens (words).
- For the words a, and, the, i, is compute their respective counts as the book progresses, i.e.

$$n_{\text{the}}[i] = \sum_{j=1}^i \{w_j = \text{the}\}$$

- Plot the proportions  $n_{\text{word}}[i]/i$  over the document in one plot.
- Find an envelope of the shape  $O(1/\sqrt{i})$  for each of these five words. (Hint, check the last page of the [sampling notebook](http://courses.d2l.ai/berkeley-stat-157/slides/1_24/sampling.pdf) ([http://courses.d2l.ai/berkeley-stat-157/slides/1\\_24/sampling.pdf](http://courses.d2l.ai/berkeley-stat-157/slides/1_24/sampling.pdf)))
- Why can we **not** apply the Central Limit Theorem directly?
- How would we have to change the text for it to apply?

- Why does it still work quite well?

**2.1 Answer)** Download <https://www.gutenberg.org/ebooks/84> (<https://www.gutenberg.org/files/84/84-0.txt>) from Project Gutenberg

- Remove punctuation, uppercase / lowercase, and split the text up into individual tokens (words).
- For the words a, and, the, i, is compute their respective counts as the book progresses, i.e.

$$n_{\text{the}}[i] = \sum_{j=1}^i \{w_j = \text{the}\}$$

- Plot the proportions  $n_{\text{word}}[i]/i$  over the document in one plot.

```

In [20]: filename = gluon.utils.download('https://www.gutenberg.org/files/84/84-0.txt')
with open(filename) as f:
    book = f.read()

def preprocess(text):
    punctuation = set(['.', '-', '(', ')', ';', '/', '*', '"', ' ', ':', "'", '['])
    nums = set('0123456789$')

    return [word.lower() for word in text.split() if not any((x in word) for x in punctuation) and not a

lst = preprocess(book)

a_counts = np.zeros(len(lst))
and_counts = np.zeros(len(lst))
the_counts = np.zeros(len(lst))
i_counts = np.zeros(len(lst))
is_counts = np.zeros(len(lst))

for i in range(0, len(lst)):
    if (i != 0):
        a_counts[i] = a_counts[i-1]
        and_counts[i] = and_counts[i-1]
        the_counts[i] = the_counts[i-1]
        i_counts[i] = i_counts[i-1]
        is_counts[i] = is_counts[i-1]
        if (lst[i] == 'a'):
            a_counts[i] += 1
        elif (lst[i] == 'and'):
            and_counts[i] += 1
        elif (lst[i] == 'the'):
            the_counts[i] += 1
        elif (lst[i] == 'i'):
            i_counts[i] += 1
        elif (lst[i] == 'is'):
            is_counts[i] += 1
    else:
        if (lst[0] == 'a'):
            a_counts[0] = 1
        elif (lst[0] == 'and'):
            and_counts[0] = 1
        elif (lst[0] == 'the'):
            the_counts[0] = 1

```

```

        elif (lst[0] == 'i'):
            i_counts[0] = 1
        elif (lst[0] == 'is'):
            is_counts[0] = 1

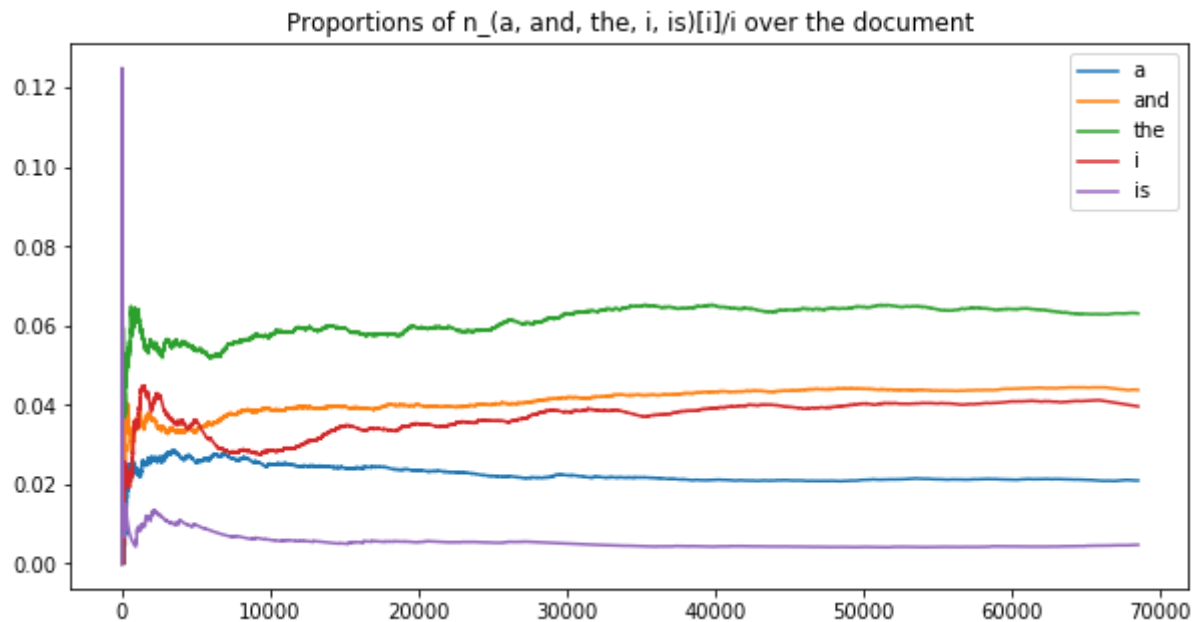
a_prop = np.zeros(len(lst))
and_prop = np.zeros(len(lst))
the_prop = np.zeros(len(lst))
i_prop = np.zeros(len(lst))
is_prop = np.zeros(len(lst))

for i in range(1, len(lst)):
    a_prop[i] = a_counts[i]/i
    and_prop[i] = and_counts[i]/i
    the_prop[i] = the_counts[i]/i
    i_prop[i] = i_counts[i]/i
    is_prop[i] = is_counts[i]/i

```

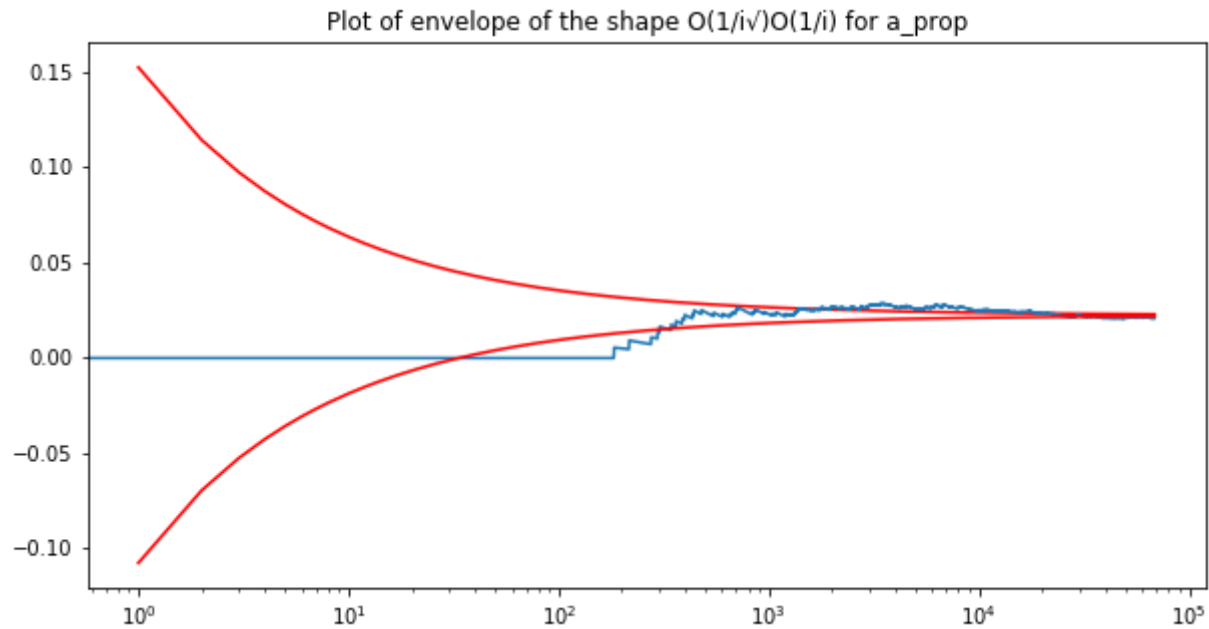
**2.2 Answer) Plot the proportions  $n_{\text{word}}[i]/i$  over the document in one plot.**

```
In [147]: plt.figure(figsize=(10,5))
plt.plot(a_prop, label='a')
plt.plot(and_prop, label='and')
plt.plot(the_prop, label='the')
plt.plot(i_prop, label='i')
plt.plot(is_prop, label='is')
plt.title('Proportions of n_(a, and, the, i, is)[i]/i over the document')
plt.legend()
plt.show()
```



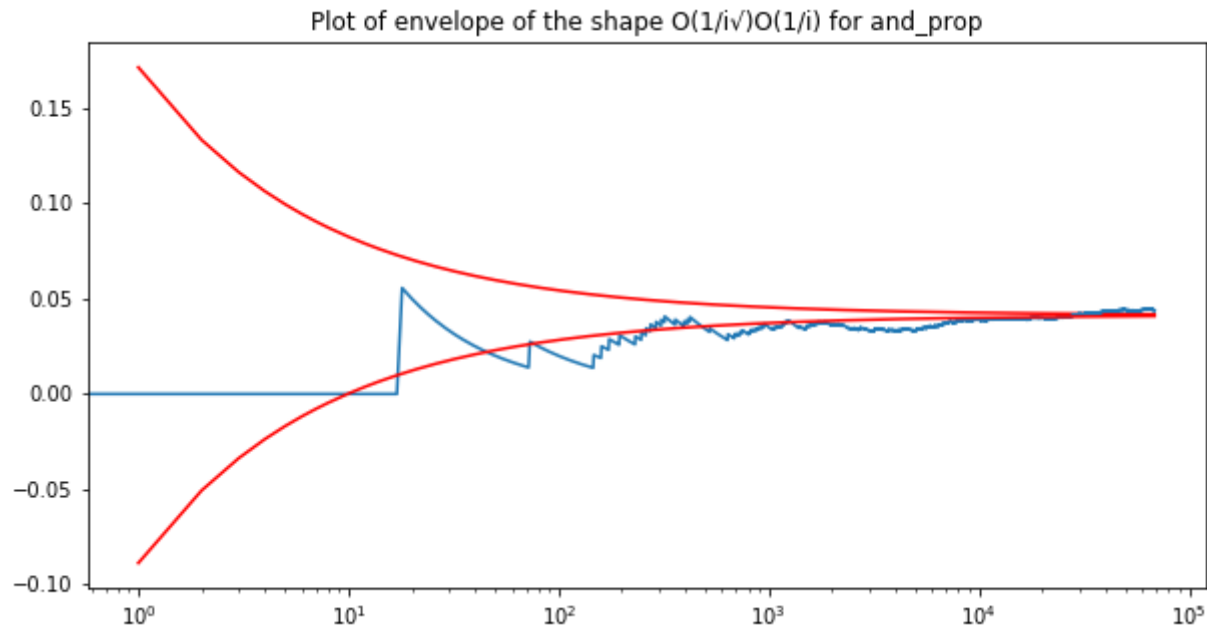
**2.3. Answer) Plot of envelope of the shape  $O(1/i_\sqrt{ })O(1/i)$  for a\_prop.**

```
In [148]: y = np.arange(1, len(lst)+1).reshape(len(lst), 1)
plt.figure(figsize=(10, 5))
plt.semilogx(y, 0.13*np.power(y, -0.5) + np.mean(a_prop), 'r')
plt.plot(a_prop)
plt.semilogx(y, -0.13*np.power(y, -0.5) + np.mean(a_prop), 'r')
plt.title('Plot of envelope of the shape  $O(1/\sqrt{i})O(1/i)$  for a_prop')
plt.show()
```



**2.3. Answer) Plot of envelope of the shape  $O(1/\sqrt{i})O(1/i)$  for and\_prop.**

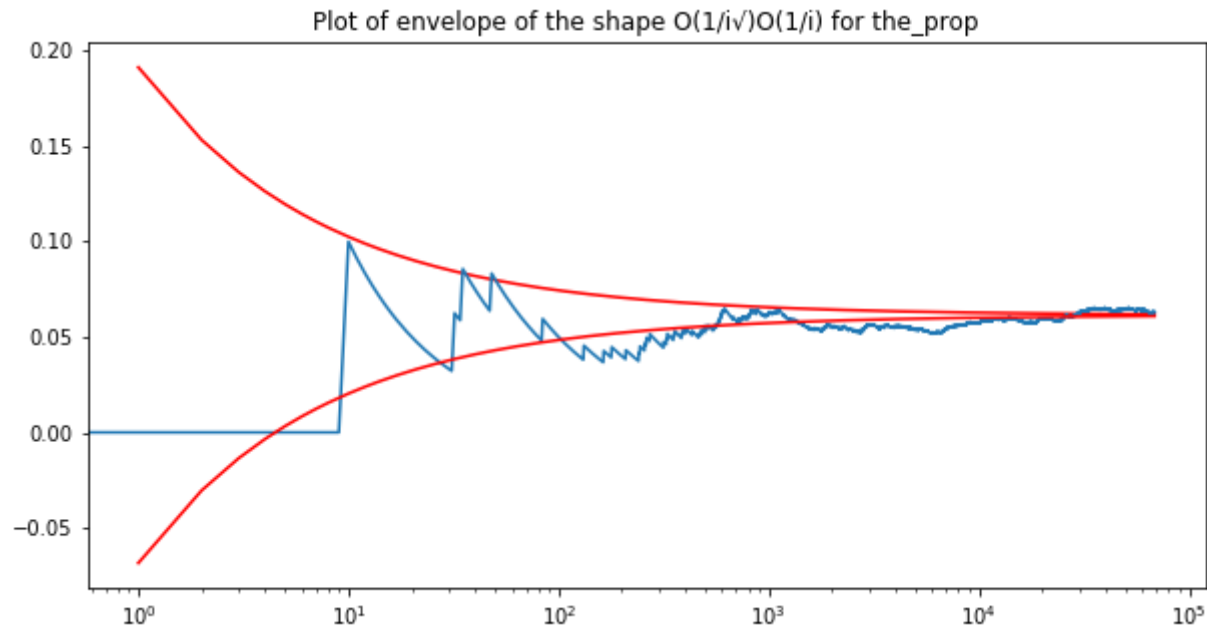
```
In [149]: y = np.arange(1,len(lst)+1).reshape(len(lst),1)
plt.figure(figsize=(10,5))
plt.semilogx(y, 0.13*np.power(y,-0.5) + np.mean(and_prop), 'r')
plt.plot(and_prop)
plt.semilogx(y, -0.13*np.power(y,-0.5) + np.mean(and_prop), 'r')
plt.title('Plot of envelope of the shape  $O(1/\sqrt{i})O(1/i)$  for and_prop')
plt.show()
```



**2.3. Answer) Plot of envelope of the shape  $O(1/\sqrt{i})O(1/i)$  for the\_prop.**

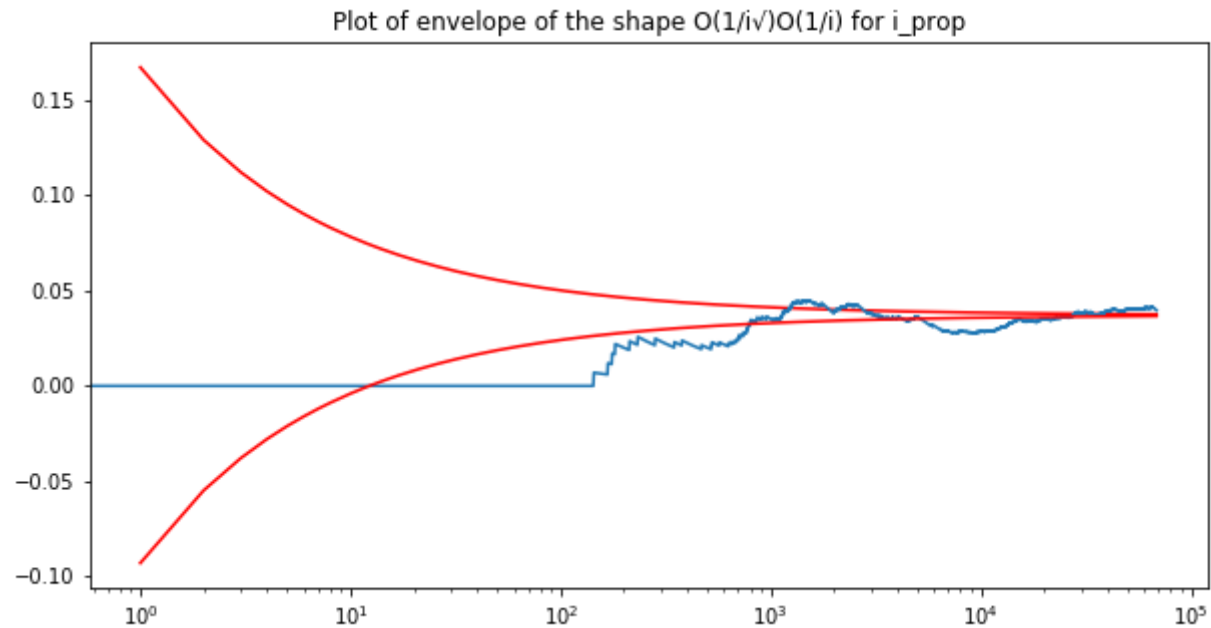


```
In [150]: y = np.arange(1,len(lst)+1).reshape(len(lst),1)
plt.figure(figsize=(10,5))
plt.semilogx(y, 0.13*np.power(y,-0.5) + np.mean(the_prop), 'r')
plt.plot(the_prop)
plt.semilogx(y, -0.13*np.power(y,-0.5) + np.mean(the_prop), 'r')
plt.title('Plot of envelope of the shape  $O(1/\sqrt{i})O(1/i)$  for the_prop')
plt.show()
```



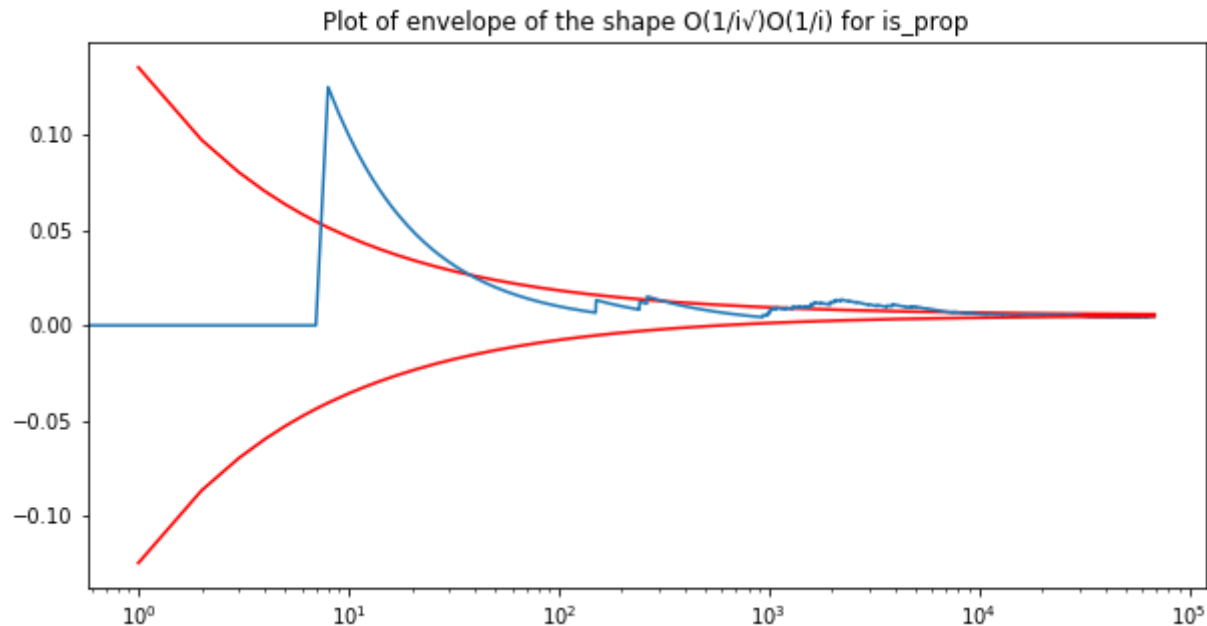
**2.3. Answer) Plot of envelope of the shape  $O(1/\sqrt{i})O(1/i)$  for  $i_{\text{prop}}$ .**

```
In [151]: y = np.arange(1,len(lst)+1).reshape(len(lst),1)
plt.figure(figsize=(10,5))
plt.semilogx(y, 0.13*np.power(y,-0.5) + np.mean(i_prop), 'r')
plt.plot(i_prop)
plt.semilogx(y, -0.13*np.power(y,-0.5) + np.mean(i_prop), 'r')
plt.title('Plot of envelope of the shape  $O(1/\sqrt{i})O(1/i)$  for  $i_{\text{prop}}$ ')
plt.show()
```



**2.3. Answer) Plot of envelope of the shape  $O(1/\sqrt{i})O(1/i)$  for  $i_{\text{prop}}$ .**

```
In [152]: y = np.arange(1,len(lst)+1).reshape(len(lst),1)
plt.figure(figsize=(10,5))
plt.semilogx(y, 0.13*np.power(y,-0.5) + np.mean(is_prop), 'r')
plt.plot(is_prop)
plt.semilogx(y, -0.13*np.power(y,-0.5) + np.mean(is_prop), 'r')
plt.title('Plot of envelope of the shape  $O(1/i^{\sqrt{}})O(1/i)$  for is_prop')
plt.show()
```



**Answer) Why can we not apply the Central Limit Theorem directly?**

We can't apply the Central Limit Theorem directly because the sequence in which words appear in a document is not independent of one another.

**Answer) How would we have to change the text for it to apply?**

If we shuffle the texts, then we can apply the CLT.

**Answer) Why does it still work quite well?**

It still works well because sample size (# of words) is large and our bag of words are relatively small so dependency isn't magnified.

### 3. Denominator-layout notation

We used the numerator-layout notation for matrix calculus in class, now let's examine the denominator-layout notation.

Given  $x, y \in \mathbb{R}$ ,  $\mathbf{x} \in \mathbb{R}^n$  and  $\mathbf{y} \in \mathbb{R}^m$ , we have

$$\frac{\partial y}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y}{\partial x_1} \\ \frac{\partial y}{\partial x_2} \\ \vdots \\ \frac{\partial y}{\partial x_n} \end{bmatrix}, \quad \frac{\partial \mathbf{y}}{\partial x} = \left[ \frac{\partial y_1}{\partial x}, \frac{\partial y_2}{\partial x}, \dots, \frac{\partial y_m}{\partial x} \right]$$

and

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial \mathbf{y}}{\partial x_1} \\ \frac{\partial \mathbf{y}}{\partial x_2} \\ \vdots \\ \frac{\partial \mathbf{y}}{\partial x_n} \end{bmatrix} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1}, \frac{\partial y_2}{\partial x_1}, \dots, \frac{\partial y_m}{\partial x_1} \\ \frac{\partial y_1}{\partial x_2}, \frac{\partial y_2}{\partial x_2}, \dots, \frac{\partial y_m}{\partial x_2} \\ \vdots \\ \frac{\partial y_1}{\partial x_n}, \frac{\partial y_2}{\partial x_n}, \dots, \frac{\partial y_m}{\partial x_n} \end{bmatrix}$$

Questions:

1. Assume  $\mathbf{y} = f(\mathbf{u})$  and  $\mathbf{u} = g(\mathbf{x})$ , write down the chain rule for  $\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$
2. Given  $\mathbf{X} \in \mathbb{R}^{m \times n}$ ,  $\mathbf{w} \in \mathbb{R}^n$ ,  $\mathbf{y} \in \mathbb{R}^m$ , assume  $z = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$ , compute  $\frac{\partial z}{\partial \mathbf{w}}$ .

Answers:

1.  $\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \frac{\partial \mathbf{u}}{\partial \mathbf{x}} \frac{\partial \mathbf{y}}{\partial \mathbf{u}}$  (performed in this order for matrix dimensions to match denominator layout notation.)

2. First expand  $z$ .  $z = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 = (\mathbf{X}\mathbf{w} - \mathbf{y})^T(\mathbf{X}\mathbf{w} - \mathbf{y})$

$$= (\mathbf{w}^T \mathbf{X}^T - \mathbf{y}^T)(\mathbf{X}\mathbf{w} - \mathbf{y})$$

$$= \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y}$$

$$= \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y}$$

Now  $\frac{\partial z}{\partial \mathbf{w}}(\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w} - 2\mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{y}^T \mathbf{y}) = 2\mathbf{w}^T (\mathbf{X}^T \mathbf{X}) - 2\mathbf{y}^T \mathbf{X}$

## 4. Numerical Precision

Given scalars  $x$  and  $y$ , implement the following `log_exp` function such that it returns

$$-\log\left(\frac{e^x}{e^x + e^y}\right)$$

```
In [145]: def log_exp(x, y):
          ## add your solution here
          return -nd.log(nd.exp(x) / (nd.exp(x) + nd.exp(y)))
```

Test your codes with normal inputs:

```
In [146]: x, y = nd.array([2]), nd.array([3])
          z = log_exp(x, y)
          z
```

```
Out[146]: [1.3132617]
          <NDArray 1 @cpu(0)>
```

Now implement a function to compute  $\partial z / \partial x$  and  $\partial z / \partial y$  with autograd

```
In [138]: def grad(forward_func, x, y):  
    ## Add your codes here  
    x.attach_grad()  
    y.attach_grad()  
    with autograd.record():  
        z = forward_func(x,y)  
    z.backward()
```

Test your codes, it should print the results nicely.

```
In [139]: grad(log_exp, x, y)  
print('x.grad =', x.grad)  
print('y.grad =', y.grad)
```

```
x.grad =  
[-0.7310586]  
<NDArray 1 @cpu(0)>  
y.grad =  
[0.7310586]  
<NDArray 1 @cpu(0)>
```

But now let's try some "hard" inputs

```
In [140]: x, y = nd.array([50]), nd.array([100])  
grad(log_exp, x, y)  
print('x.grad =', x.grad)  
print('y.grad =', y.grad)
```

```
x.grad =  
[nan]  
<NDArray 1 @cpu(0)>  
y.grad =  
[nan]  
<NDArray 1 @cpu(0)>
```

Does your code return correct results? If not, try to understand the reason. (Hint, evaluate  $\exp(100)$ ). Now develop a new function `stable_log_exp` that is identical to `log_exp` in math, but returns a more numerical stable result.

```
In [143]: def stable_log_exp(x, y):
            if x > y:
                a = x
                b = y
            else:
                a = y
                b = x
            return a + nd.log(1 + nd.exp(b-a)) - x
grad(stable_log_exp, x, y)
print('x.grad =', x.grad)
print('y.grad =', y.grad)

x.grad =
[-1.]
<NDArray 1 @cpu(0)>
y.grad =
[1.]
<NDArray 1 @cpu(0)>
```

```
In [ ]:
```